

Gaurav Kumar

New York, NY | [linkedin.com/in/gaurav98](https://www.linkedin.com/in/gaurav98) | gaurav.kumar9825@gmail.com | github.com/liquidslr

Skills

Languages: Java, C++, Python, JavaScript, Golang

Databases: MySQL, DynamoDB, PostgreSQL, MongoDB, Redis

Generative AI: LLMs, RAG, Fine-tuning (LoRA/QLoRA), Diffusion Models

Cloud Platforms: AWS (EC2, S3, SQS, Lambda, CloudWatch, ECS)

Tools: Docker, Kubernetes, Terraform, LangChain

Work Experience

Amazon Web Services (AWS)

New York, NY

Software Engineer

07/2025 – Present

- **Engineered** the Contract Data Model (CDM) Library, a versioned, strongly-typed clause-level model with auto-generated extraction prompts and snapshot-based regression tests, now adopted by **four** upstream teams; cut new clause-variant development to **under 20 minutes** and enabled auto-ingestion to scale from **1.4%** toward an **80%** target across **~40,000** contracts.
- **Owned** the preceding hybrid rule-based/LLM extraction engine that pulls pricing terms from executed AWS enterprise agreements into AWS billing (**~1,200** contracts/month, **14,000+** to date), which CDM now replaces; hardened reliability with idempotency, retries, and **P50/P90/P99** latency instrumentation.
- **Automated** end-to-end CloudFront Opt-out contract termination via an event-driven pipeline with retry, dedup, and reversal handling, eliminating a manual monthly scrub of **~100** agreements and raising on-time onboarding from **60%** to **95.7%**.

Yojak

Gurgaon, India

Software Engineer

05/2022 – 12/2023

- **Built** a React Native cross-platform e-commerce app (QR SKU scanning, real-time stock, fault-tolerant transactions) that lifted user engagement by **35%**, and architected an integrated ERP consolidating order processing, accounting, and supply chain, automating **15+** manual processes to save **40+ hours/week**.

Fractal Analytics

Mumbai, India

Data Scientist

10/2020 – 05/2022

- **Built** a Customer Lifetime Value model to optimize retail marketing spend across channels, driving a **\$10M** annual revenue lift, and an implicit collaborative filtering recommender with seasonal prioritization that lifted sales **7%** (**\$100k/month**), validated via A/B testing.
- **Increased** lead conversion by **20%** with a tree-based lead scoring model deployed for real-time inference on Kubernetes with FastAPI.

Education

University of South Carolina, Molinaroli College of Engineering and Computing

Columbia, SC

Master's in Computer Science — GPA: 4.0 / 4.0

01/2024 – 05/2025

Indian Institute of Technology (IIT), Roorkee

Uttarakhand, India

Bachelor's in Engineering Physics

07/2016 – 05/2020

Projects

DragAI

01/2025

- Built an open-source RAG platform that crawls and indexes entire websites into vector and relational stores, serving grounded LLM answers over the extracted content; designed the chunking and retrieval pipeline to handle multi-URL ingestion and incremental re-indexing.

Publications

- Rawte, Scott, [Kumar](#), et al. “Do Voters Get the Information They Want? Understanding Authentic Voter FAQs in the US.” *TrustNLP Workshop, NAACL 2025*.
- Jajee, [Kumar](#), Valafar. “Automated Vigilance State Classification in Rodents Using Machine Learning.” *arXiv:2507.14166* (2025).

Awards

5th Big Data Health Science Student Case Competition

01/2024

- Winner out of 30 teams for a sleep-stage classification model.